

Python for Data processing

Lecture 3:

Arrays, tensors and computations - Part III

Kosta Rozen

What we already know

A lot about NumPy arrays:

- creation
- indexing
- universal functions
- linear algebra
- best practices
- I/O

This lecture

PyTorch:

- basics
- operations
- gradients
- logistic regression

Why NumPy is not enough

NumPy arrays are great but:

- they work only on CPU
- they provide only basic building blocks

For deep learning:

- CPU/GPU/TPU/?
- gradients

PyTorch

- **tensors** provide the same operations as NumPy arrays
- work on CPU/**GPU/TPU**
- provide **autogradients**
- **deep learning** building blocks
- efficient **data loading**
- **deployment**


PyTorch tensors

Tensors


- similar to arrays, provide the same computational facilities
- can **share** data
- can live on **different devices**
- provide **declarative** computations

→let's try it out!

Tensors and storage



Devices and computations



View, copies, reshaping

PyTorch is a bit more elaborated:

- **view**: always returns a view or fails
- **reshape**: returns either view or new tensor
- depends on contiguity constraints

Gradients

In deep learning we need **gradients**:

- to calculate updates to network parameters (weights)
- no way to do that in NumPy
- an easy go in PyTorch (a bit more elaborated in Tensorflow)

$$L(a_{ij}) \text{ (scalar)} \rightarrow \frac{\partial L}{\partial a_{ij}} \text{ (tensor)}$$




→let's try it out!

Logistic regression with PyTorch

Logistic regression: setup

- **two-dimensional** input (created with `make_blobs` from `sklearn.datasets`)
- **binary** classification with **linear** decision boundary
- output: **sigmoid**
- from **scratch**

LR breakdown



Scalar product of W and x : $Wx = W_0x_0 + W_1x_1$


Add bias b : $W_0x_0 + W_1x_1 + b$



$$P(\bullet) = \sigma(Wx+b)$$



Apply Sigmoid function:
If $Wx+b > 0$, then $P(\bullet) > 0.5$
If $Wx+b < 0$, then $P(\bullet) < 0.5$



$P(\bullet) = P(\bullet) - \text{decision boundary}$

Log loss

class labels



sigmoid
output



good

bad

We want: probability ↓ for class 0, probability ↑ for class 1

Log loss function

Log loss is $-1 \times$ the log-likelihood function you learned in probability course as part of max-likelihood method for parameter estimation:

The diagram shows the log loss function formula with several annotations:

- True value of sample k:** y_k (1 for red or 0 for blue)
- Predicted probability of sample k being red:** \hat{y}_k
- Predicted probability of sample k being blue:** $1 - \hat{y}_k$
- Number of samples:** N
- 1 if sample k is blue:** $1 - y_k$

$$L = -\frac{1}{N} \sum_k (y_k \log(\hat{y}_k) + (1 - y_k) \log(1 - \hat{y}_k))$$

The formula is annotated with arrows pointing to the following components:

- N : Number of samples
- y_k : True value of sample k: 1 for red or 0 for blue
- \hat{y}_k : Predicted probability of sample k being red
- $1 - \hat{y}_k$: Predicted probability of sample k being blue
- $1 - y_k$: 1 if sample k is blue

The sigmoid function is defined as:

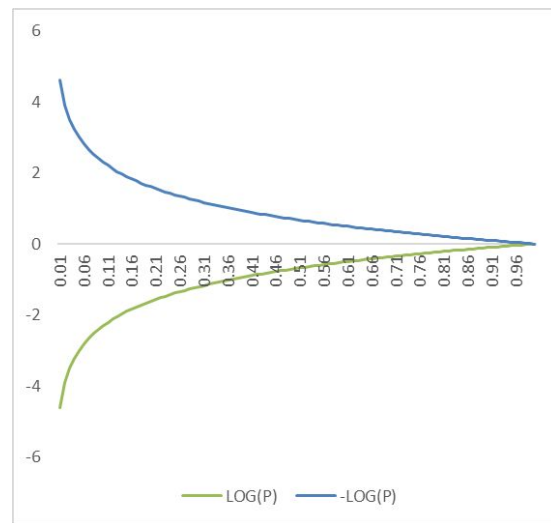
$$\hat{y}_i = \sigma(X_{ik} W_k + b)$$

Log loss function

Log parts of the function represent the cost of error when estimating probability:

- When we're right ($p=1$), $\log(p)$ is 0
- When we're wrong (small p), $\log(p)$ is negative and $-\log(p)$ is positive.
- The smaller p is, the larger the error ($-\log$) will be.

$$L = -\frac{1}{N} \sum_k (y_k \log(\hat{y}_k) + (1 - y_k) \log(1 - \hat{y}_k))$$
$$\hat{y}_i = \sigma(X_{ik}W_k + b)$$



→let's try it out!

What we've learned

- PyTorch tensors and gradients
- how to perform simple gradient descent

Assignment

- explore PyTorch tensor operations

questions?