

# Python for Data processing

## Lecture 4: **Pandas - Part I**

Kosta Rozen

# What we already know

**A lot about NumPy arrays and PyTorch tensors:**

- basic operations
- best practices
- optimization
- linear algebra
- very basics of how machine learning works

# Structured and unstructured data

# Unstructured data

- images
- signals (including time series)
- text

Each data element (pixel, datapoint, letter) is usually atomic and **is equal** to any other data element. You need to **perform analysis** to get the structure

# Structured data

- tabular data
- JSON
- XML

Each data element (row, DB record, XML file) has **internal structure** or **schema**

```
[{'name': 'Anny Smith', 'age': 35, 'sex': 'female'},  
 {'name': 'John Black', 'age': 62, 'sex': 'male'}, ...]
```

# Dataframe

**Tabular representation** of structured data

- well known in R world for years
- **indexed** rows and columns
- SQL-like operations<sup>(joins, filtering)</sup>, aggregations, alignment and more

# Pandas

One of the most respected Python packages for data science

- started in **2008**
- **very fast** (a lot of Cython inside)
- supports **tons of operations and formats**
- extremely **flexible** and **powerful**
- It's **crazy** sometimes, but you'll **love** it


# Pandas series and df's

Pandas has two main data structures:

- **pd.Series** for indexed 1D data
- **pd.DataFrame** for indexed tabular data




# Pandas series




→let's try it out!

# Pandas dataframe



# Pandas dataframe



# Indexing series and dataframes

# Reading csv files

The best tool to read CSV and other text files in Python:

```
pd.read_csv(...)
```

→let's try it out!

# Indexing series and df's

- `[]` indexing
- `.loc` - label based indexing
- `.iloc` - position based indexing

**Boolean indexing** is possible and is heavily used

→let's try it out!

# Indexing df: SettingWithCopyWarning

Pandas is not like numpy:

- **It's unknown whether you get view or copy**

Why?

- **It's hard to give guarantees**  
(but there are rules inside) (but you should not even try to understand them)

→let's try it out!

# Operations on dataframes




# Arithmetic: not your usual NumPy

Pandas **aligns** dataframes for you before performing operations by creating a **union** of row and column indexes

→let's try it out!

# Arithmetic



→let's try it out!

# Applying functions to df's

Pandas allows you to apply custom functions across rows or columns, and elementwise.

And combines results for you appropriately.

It's usually **fast**.

→let's try it out!

# Dataframe summaries

It's easy to get general information about dataframe:

`df.info()`

`df.describe()`

`df.head()`, `df.tail()`

→let's try it out!

# Counts and statistics

To get counts or statistics about column or row:

```
df[col].unique()
```

```
df[col].value_counts()
```

```
df.sum(axis=...), df.mean(axis=...), df.std(axis=...)
```

Powerful in combination with smart indexing.

→let's try it out!

# Replacing and renaming

**df.replace** allows for flexible replacement of values in dataframes:

- by value, per column

**df.rename** allows you to easily rename any label, be it column name or index label

→let's try it out!

# Missing data

Pandas is great at handling missing data:

- infers it for you
- backward fill, forward fill and more

→let's try it out!

# Categorical data

Pandas easily calculates one-hot encoded values for any column, adding properly named columns

→let's try it out!



# Special datatypes

Pandas has very good support for:

- **strings** - great for text columns  
(split, replace and other usual string operations, vectorized)
- **datetime** - flexible indexing, handling timezones and extravagant parsing  
(great for anything time series related)

→let's try it out!

# What we've learned

Basics of Pandas:

- creating, indexing
- operations on dataframes

# Assignment

- Explore Pandas
- play with Titanic dataset
- 2 weeks!

questions?